

# Machine Learning Pipelines

---

## ➤ *What is a Machine Learning Pipeline?*

A **pipeline** in machine learning is a step-by-step *assembly line* that automatically handles everything from **data preprocessing** to **model training**—all in a single streamlined structure.

---

## ➤ *What Does It Include?*

A typical ML pipeline includes:

1. **Data cleaning & preprocessing** (e.g., handling missing values)
2. **Feature transformation** (e.g., scaling, encoding)
3. **Model training**
4. **Evaluation**
5. **Model saving (Pickling)**
6. **Deployment or Prediction**

All of this is tied together so you don't have to manually repeat steps.

---

## ➤ *Why Use Pipelines?*

- **Clean code:** Everything is organized in one structure
  - **Reproducibility:** You can rerun the same process reliably
  - **Less error-prone:** Avoid mistakes from copy-pasting code
  - **Works with cross-validation & grid search easily**
  - **Production ready:** You can save and load pipelines directly
- 

## ➤ *Example(Titanic Dataset):*

In that dataset, you went through:

- Splitting the Titanic dataset
  - Imputing missing values (e.g., age)
  - Creating a pipeline with preprocessing + model
  - Comparing Pipeline vs make\_pipeline (both create pipelines, but Pipeline lets you name the steps)
  - Pickling (saving) the model
  - Cross-validating using the full pipeline
-

➤ **Pipeline vs make\_pipeline:**

- Pipeline: You define and name each step explicitly (good for customization).
  - make\_pipeline: Quick and automatic, but steps are unnamed.
- 

➤ **What is Pickling?**

Pickling is just saving your trained pipeline (model + preprocessing steps) so that you can reuse it later—without retraining from scratch.

---

➤ **Summary:**

**A Machine Learning Pipeline is your complete ML workflow automated, efficient, and ready to deploy.**

It's like wrapping your entire process in one neat box and pressing *play*.

---