

KNN Imputer & Multivariate Imputation

➤ Handling Missing Data

This lecture covers **KNN Imputer**, a **multivariate** method of imputing missing data using patterns across multiple variables.

1. Univariate vs. Multivariate Imputation:

- **Univariate:** Fills missing values using information from the **same column only** (e.g., mean, median).
 - **Multivariate:** Considers **multiple columns** to infer missing values — more context-aware and often more accurate.
-

2. How KNN Imputer Works:

- Finds the **k-nearest neighbors** (based on Euclidean distance) for the row with missing data.
 - Uses **values from those neighbors** to estimate the missing value (mean by default).
 - **Distance is calculated using available (non-missing) features only** — see scikit-learn's documentation.
-

3. Euclidean Distance with NaNs:

- The lecture explains how **distance is calculated when some features are missing**.
- scikit-learn handles this using a **masked Euclidean distance**, ignoring NaNs during the calculation.

Documentation:

- [KNN Imputer - scikit-learn](#)
-

4. Pros & Cons of KNN Imputer

❖ Advantages:

- Preserves **relationships between features**.
- No assumption about data distribution.

❖ Disadvantages:

- **Computationally expensive** with large datasets.
 - **Sensitive to outliers** and irrelevant features.
 - Needs **scaling** for accurate distance computation.
-

5. Uniform vs. Distance Weighting:

- `weights='uniform'`: All neighbors contribute **equally**.
-

- `weights='distance'`: **Closer neighbors** contribute more to the imputed value.

Use `weights='distance'` when you want **more precision** from closer data points.

➤ **Key Takeaway:**

KNN Imputer is a **smart multivariate technique** that uses neighbor similarity to fill in missing values. It's **more powerful than simple methods** but needs careful tuning (especially for large or unscaled data).
