

Logistic Regression Gradient Descent

1. Gradient Descent

- **What it is:**

An optimization algorithm used to find the parameters (weights) that minimize the loss function.

- **Idea:**

1. Start with random weights.
2. Calculate the loss (error) between predictions and actual labels.
3. Update weights in the opposite direction of the gradient (slope) to reduce error.

- **Update Rule:**

$$w_{\text{new}} = w_{\text{old}} - \alpha \cdot \frac{\partial L}{\partial w}$$

where α is the learning rate.

- **In Logistic Regression:**

The gradient is computed from the derivative of the **Binary Cross-Entropy loss** with respect to weights, which depends on the difference between predicted probabilities and actual labels.

2. Loss Function in Matrix Form

- **Why matrix form?**

1. Makes calculations efficient for multiple data points using vectorized operations (fast, less error-prone).

- **Matrix Loss Formula:**

For m samples:

$$L = -\frac{1}{m} [y^T \log(\hat{y}) + (1 - y)^T \log(1 - \hat{y})]$$

where:

1. y = vector of actual labels (0 or 1)
2. \hat{y} = vector of predicted probabilities from the sigmoid function.

3. Code Demo

Typical steps when coding Logistic Regression from scratch:

1. **Initialize** weights and bias to zeros (or small random values).
2. **Forward pass:**

1. Compute $z = Xw + b$
3. Apply sigmoid: $\hat{y} = \sigma(z)$
4. **Loss calculation:** Using Binary Cross-Entropy in matrix form.
5. **Backward pass:**
 1. Compute gradient for weights:

$$\frac{\partial L}{\partial w} = \frac{1}{m} X^T (\hat{y} - y)$$

2. Compute gradient for bias:

$$\frac{\partial L}{\partial b} = \frac{1}{m} \sum (\hat{y} - y)$$

6. **Parameter update:** Apply gradient descent update rule for w and b .
7. **Repeat** for several iterations until loss converges.

➤ **Key takeaway:**

Gradient Descent iteratively tweaks weights to make predictions closer to the truth. In Logistic Regression, this process uses the sigmoid output, Binary Cross-Entropy loss, and vectorized math for speed.
